

# Rexample3


Lingxiao Zhou

2024-09-16

## Example for ANOVA (toluca)

Data on lot size ( $x$ ) and work hours ( $y$ ) was obtained from 25 recent runs of a manufacturing process.

```
# find relationship between lot size (X) and work hours (Y)
toluca = read.table("http://www.stat.ufl.edu/~rrandles/sta4210/Rclassnotes/data/textdatasets/KutnerData.txt",
  col.names = c("lotsize", "workhrs"))
plot(toluca$lotsize, toluca$workhrs, xlab = "lot size", ylab = "work hours")
toluca.reg = lm(workhrs ~ lotsize, data = toluca)
abline(toluca.reg)
```



```
anova(toluca.reg)
```

```
## Analysis of Variance Table
##
## Response: workhrs
##          Df Sum Sq Mean Sq F value    Pr(>F)
## lotsize   1 252378  252378   105.88 4.449e-10 ***
```

```
## Residuals 23 54825 2384
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(toluca.reg)
```

```
##
## Call:
## lm(formula = workhrs ~ lotsize, data = toluca)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -83.876 -34.088  -5.982  38.826 103.528
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   62.366     26.177   2.382  0.0259 *
## lotsize        3.570       0.347  10.290 4.45e-10 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 48.82 on 23 degrees of freedom
## Multiple R-squared:  0.8215, Adjusted R-squared:  0.8138
## F-statistic: 105.9 on 1 and 23 DF,  p-value: 4.449e-10
```

**Reduce Model:**  $Y_i = \beta_0$

**Full Model:**  $Y_i = \beta_0 + \beta_1 x_i$

```
Red <- lm(workhrs ~ 1, data = toluca)
Full <- toluca.reg
```

```
anova(Red, Full)
```

```
## Analysis of Variance Table
##
## Model 1: workhrs ~ 1
## Model 2: workhrs ~ lotsize
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1      24 307203
## 2      23  54825  1   252378 105.88 4.449e-10 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


## Diagnostic on Predictor

**Goal:** Identify any outlying values in predictor,  $x$  that could affect the appropriateness of the linear model.

**Two main issues**


- Outliers (Histogram and/or Boxplot)
- Levels of predictor associated with the run order when experiment is run sequentially (Sequence/ Time Series Plot)

```
par(mfrow = c(1, 2))
hist(toluca$lotsize, xlab = "lot size", main = "Histogram")
boxplot(toluca$lotsize)
```



- Do not appear to be any outliers

```
plot(toluca$lotsize, type = "b", xlab = "Run order", ylab = "Lot size",
     main = "Sequence Plot")
```



- No discernible pattern/dependency of the values and run order.

## Checking Model Assumptions Graphically

Inference are only valid if the model assumptions are valid.

$$Y_i = \beta_0 + \beta_1 x_i + \epsilon_i, \quad i = 1, \dots, n,$$

where  $\epsilon_i \stackrel{iid}{\sim} N(0, \sigma)$  for  $i = 1, \dots, n$ .

- Normality: The errors are normally distributed
- Independence: The errors are independent of each other.
- Homogeneity of Variance: The variance of the residuals is the same for all values of X.
- Linearity: The relationship between the independent variable ( $x$ ) and the mean of the dependent variable ( $Y$ ) is linear.

## Residuals

- Assumptions can be checked with residuals  $e_i := y_i - \hat{y}_i$ , or better yet, the standardized/studentized residuals
- Standardized residuals have mean 0 and standard deviation of 1
- Studentized residuals are standardized residuals that use the leave-one-out approach. The  $i$ -th observation is omitted when fitting the model is then used to predict (without bias) the response for the  $i$ -th observation. More on this later
- In R, use `rstandard()` and `rstudent()` function to obtain the standardized and studentized residual.

# Normality

## Empirical CDF

**Def:** CDF associated with the smprical measuare of the sample. Assigns equal probability  $1/n$  to each point and is a step function

$$\hat{F}_n(x) = \frac{\text{number of elements} \leq x}{n} = \frac{1}{n} \sum_{i=1}^n I\{x_i \leq x\}$$


where  $I(\cdot)$  is the indicator function.

**Example:** consider the sample 1, 3, 7, 8

$$\hat{F}_4(x) = \begin{cases} 0 & \text{if } x < 1 \\ 0.25 & \text{if } 1 \leq x < 3 \\ 0.5 & \text{if } 3 \leq x < 7 \\ 0.75 & \text{if } 7 \leq x < 8 \\ 1 & \text{if } x \geq 8 \end{cases}$$

```
sample_data <- c(1, 3, 7, 8)
edf <- ecdf(sample_data)

# Plot the empirical distribution function
plot(edf, main = "Empirical CDF", xlab = "x", ylab = "F(x)",
     verticals = FALSE, do.points = FALSE)
```




- The more data points the smaller the jumps and the step function begins to look like an “S” curve

**Example:** Generate 100 observations from a  $N(5,2)$  and plot the empirical CDF and overlay the theoretical CDF from  $N(5,2)$

```

set.seed(111)
y <- rnorm(100, 5, 2)
edf <- ecdf(y)
plot(edf, main = "Empirical CDF")
curve(pnorm(x, 5, 3), add = TRUE, col = "red")
legend("topleft", legend = c("ECDF", "N(5,2) CDF"), lty = 1,
      col = c("black", "red"))

```



## Q-Q plot


Let  $G(\cdot)$  denotes the theoretical normal CDF  
 Q-Q plot plots the quantile function  $\hat{F}_n^{-1}(x)$  versus  $G^{-1}(x)$

```

qqnorm(toluca.reg$residuals)
qqline(toluca.reg$residuals)

```

## Normal Q-Q Plot




## Independence

- Time series plot of residual vs time order in which it was recorded
- Sorted data may invalidate conclusion
- Independence is graphically check if there is no discernible pattern

```
plot(rstandard(toluca.reg), type = "b", xlab = "Order", ylab = "Standardized residual",  
     main = "Independence")  
abline(h = 0)
```


## Independence



## Homogeneity of variance

- Plot of the standardized/studentized residuals versus the fitted values
  - Constant variance: constant spread/distance of the residuals

```
plot(toluca.reg$fitted.values, rstandard(toluca.reg), xlab = "fitted values",  
     ylab = "standardized residuals")
```




## Linearity

- Scatter plot: For good model fit, scatterplot indicate a linear trend.
- Residual vs fitted: For good model fit, residuals should evenly spread on either side of the 0 line.

```
par(mfrow = c(1, 2))  
x <- rnorm(30, 0, 3)  
eps <- rnorm(30, 0, 2)  
y <- x^2 + eps  
plot(x, y)  
fit <- lm(y ~ x)  
abline(fit, col = "red")  
plot(fit$fitted.values, rstandard(fit), xlab = "fitted values",  
     ylab = "standardized residuals")  
abline(h = 0)
```





- We may still fit a linear model by adding or transforming predictors to include higher polynomial terms. For example,  $Y_i = \beta_0 + \beta_1 x^2 + \epsilon_i$